Continuous Integration of Machine Learning Models with ease.ml/ci: Towards a Rigorous Yet Practical Treatment

Cedric Renggli [1] **Bojan Karlaš** [1] **Bolin Ding** [2] **Feng Liu** [3] **Kevin Schawinski** [4] **Wentao Wu** [5] **Ce Zhang** [1]

ABSTRACT

Continuous integration is an indispensable step of modern software engineering practices to systematically manage the life cycles of system development. Developing a machine learning model is no difference — it is an engineering process with a life cycle, including design, implementation, tuning, testing, and deployment. However, most, if not all, existing continuous integration engines do not support machine learning as first-class citizens.

In this paper, we present ease.ml/ci, to our best knowledge, the first continuous integration system for machine learning. The challenge of building ease.ml/ci is to provide rigorous guarantees, e.g., *single accuracy point error tolerance with 0.999 reliability*, with a practical amount of labeling effort, e.g., *2K labels per test*. We design a domain specific language that allows users to specify integration conditions with reliability constraints, and develop simple novel optimizations that can lower the number of labels required by up to two orders of magnitude for test conditions popularly used in real production systems.

1 INTRODUCTION

In modern software engineering (Van Vliet et al., 2008), continuous integration (CI) is an important part of the best practice to systematically manage the life cycle of the development efforts. With a CI engine, the practice requires developers to integrate (i.e., commit) their code into a shared repository at least once a day (Duvall et al., 2007). Each commit triggers an automatic build of the code, followed by running a pre-defined test suite. The developer receives a pass/fail signal from each commit, which guarantees that every commit that receives a pass signal satisfies properties that are necessary for product deployment and/or presumed by downstream software.

Developing machine learning models is no different from developing traditional software, in the sense that it is also a full life cycle involving design, implementation, tuning, testing, and deployment. As machine learning models are used in more task-critical applications and are more tightly integrated with traditional software stacks, it becomes increasingly important for the ML development life cycle also to be managed following systematic, rigid engineering discipline. We believe that developing the theoretical and system foundation for such a life cycle management system will be an emerging topic for the SysML community.
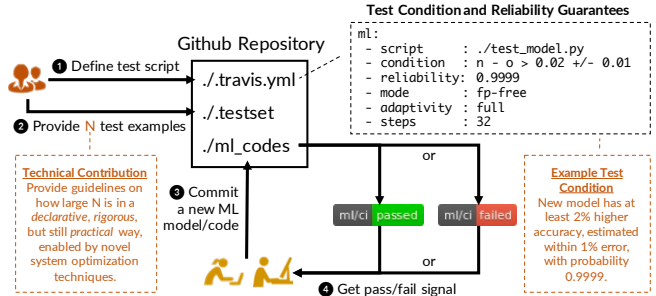

Figure 1. The workflow of ease.ml/ci.

In this paper, we take the first step towards building, to our best knowledge, the first continuous integration system for machine learning. The workflow of the system largely follows the traditional CI systems (Figure 1), while it allows the user to define machine-learning specific test conditions such as *the new model can only change at most 10% predictions of the old model* or *the new model must have at least 1% higher accuracy than the old model.* After each commit of a machine learning model/program, the system automatically tests whether these test conditions hold, and return a pass/fail signal to the developer. Unlike traditional CI, CI for machine learning is inherently *probabilistic*. As a result, all test conditions are evaluated with respect to a (ϵ, δ)-reliability requirement from the user, where $1-\delta$ (e.g., 0.9999) is the probability of a valid test and ϵ is the error tolerance (i.e., the length of the $(1-\delta)$-confidence interval). The goal of the CI engine is to return the pass/fail signal that satisfies the (ϵ, δ)-reliability requirement.

---

[1]Department of Computer Science, ETH Zurich, Switzerland [2]Data Analytics and Intelligence Lab, Alibaba Group [3]Huawei Technologies [4]Modulos AG [5]Microsoft Research. Correspondence to: Cedric Renggli <cedric.renggli@inf.ethz.ch>, Ce Zhang <ce.zhang@inf.ethz.ch>.

Technical Challenge: Practicality At the first glance of the problem, there seems to exist a trivial implementation: For each committed model, draw N labeled data points from the testset, get an (ϵ, δ)-estimate of the accuracy of the new model, and test whether it satisfies the test conditions or not. The challenge of this strategy is the practicality associated with the label complexity (i.e., how large N is). To get an $(\epsilon = 0.01, \delta = 1 - 0.9999)$ estimate of a random variable ranging in $[0, 1]$, if we simply apply Hoeffding's inequality, we need more than 46K labels from the user (similarly, 63K labels for 32 models in a non-adaptive fashion and 156K labels in a fully adaptive fashion, see Section 3)! The technical contribution of this work is a collection of techniques that lower the number of samples, by up to two orders of magnitude, that the system requires to achieve the same reliability.

In this paper, we make contributions from both the system and machine learning perspectives.

1. **System Contributions.** We propose a novel system architecture to support a new functionality compensating state-of-the-art ML systems. Specifically, rather than allowing users to compose adhoc, free-style test conditions, we design a domain specific language that is more restrictive but expressive enough to capture many test conditions of practical interest.

2. **Machine Learning Contributions.** On the machine learning side, we develop simple, but novel, optimization techniques to optimize for test conditions that can be expressed within the domain-specific language that we designed. Our techniques cover different modes of interaction (fully adaptive, non-adaptive, and hybrid), as well as many popular test conditions that industrial and academic partners found useful. For a subset of test conditions, we are able to achieve up to two orders of magnitude savings on the number of labels that the system requires.

Beyond these specific technical contributions, conceptually, this work illustrates that enforcing and monitoring an ML development life cycle in a rigorous way *does not need to be expensive*. Therefore, ML systems in the near future could afford to support more sophisticated monitoring functionality to enforce the "right behavior" from the developer.

In the rest of this paper, we start by presenting the design of ease.ml/ci in Section 2. We then develop estimation techniques that can lead to strong probabilistic guarantees using test datasets with moderate labeling effort. We present the basic implementation in Section 3 and more advanced optimizations in Section 4. We further verify the correctness and effectiveness of our estimation techniques via an experimental evaluation (Section 5). We discuss related work in Section 6 and conclude in Section 7.

2 SYSTEM DESIGN

We present the design of ease.ml/ci in this section. We start by presenting the interaction model and workflow as illustrated in Figure 1. We then present the scripting language that enables user interactions in a declarative manner. We discuss the syntax and semantics of individual elements, as well as their physical implementations and possible extensions. We end up with two system utilities, a "sample size estimator" and a "new testset alarm," the technical details of which will be explained in Sections 3 and 4.

2.1 Interaction Model

ease.ml/ci is a *continuous integration system* for machine learning. It supports a four-step workflow: (1) user describes test conditions in a *test configuration script* with respect to the quality of an ML model; (2) user provides N test examples where N is automatically calculated by the system given the configuration script; (3) whenever developer commits/checks in an updated ML model/program, the system triggers a build; and (4) the system tests whether the test condition is satisfied and returns a "pass/fail" signal to the developer. When the current testset loses its "statistical power" due to repetitive evaluation, the system also decides on when to request a new testset from the user. The old testset can then be released to the developer as a validation set used for developing new models.

We also distinguish between two teams of people: the integration team, who provides testset and sets the reliability requirement; and the development team, who commits new models. In practice, these two teams can be identical; however, we make this distinction in this paper for clarity, especially in the fully adaptive case. We call the integration team *the user* and the development team *the developer*.

2.2 A **ease.ml/ci** Script

ease.ml/ci provides a declarative way for users to specify requirements of a new machine learning model in terms of a set of test cases. ease.ml/ci then compiles such specifications into a *practical* workflow to enable evaluation of test cases with rigorous theoretical guarantees. We present the design of the ease.ml/ci scripting language, followed by its implementation as an extension to the .travis.yml format used by Travis CI.

Logical Data Model The core part of a ease.ml/ci script is a user-specified condition for the continuous integration test. In the current version, such a condition is specified over three variables $\mathcal{V} = \{n, o, d\}$: (1) n, the accuracy of the new model; (2) o, the accuracy of the old model; and (3) d, the *percentage* of new predictions that are different from the old ones ($n, o, d \in [0, 1]$).

A detailed overview over the exact syntax and its semantics is given in Appendix A.

Adaptive vs. Non-adaptive Integration A prominent difference between `ease.ml/ci` and traditional continuous integration system is that the statistical power of a test dataset will decrease when the result of whether a new model passes the continuous integration test is released to the developer. The developer, if she wishes, can adapt her next model to increase its probability to pass the test, as demonstrated by the recent work on adaptive analytics (Blum & Hardt, 2015; Dwork et al., 2015). As we will see, ensuring probabilistic guarantee in the adaptive case is more expensive as it requires a larger testset. `ease.ml/ci` allows the user to specify whether the test is adaptive or not with a flag `adaptivity (full, none, firstChange)`:

- If the flag is set to `full`, `ease.ml/ci` releases whether the new model passes the test immediately to the developer.
- If the flag is set to `none`, `ease.ml/ci` accepts all commits, however, sends the information of whether the model really passes the test to a user-specified, third-party, email address that the developer does not have access to.
- If the flag is set to `firstChange`, `ease.ml/ci` allows full adaptivity before the first time that the test passes (or fails), but stops afterwards and requires a new testset (see Section 3 for more details).

Example Scripts A `ease.ml/ci` script is implemented as an extension to the `.travis.yml` file format used in Travis CI by adding an `ml` section. For example,

```
ml:
  - script    : ./test_model.py
  - condition : n - o > 0.02 +/- 0.01
  - reliability: 0.9999
  - mode      : fp-free
  - adaptivity : full
  - steps     : 32
```

This script specifies a continuous test process that, with probability larger than 0.9999, accepts the new commit only if the new model has two points higher accuracy than the old one. This estimation is conducted with an estimation error within one accuracy point in a "false-positive free" manner. We give a detailed definition, as well as a simple example of the two modes `fp-free` and `fn-free` in Appendix A.2. The system will release the `pass/fail` signal immediately to the developer, and the user expects that the given testset can be used by as many as 32 times before a new testset has to be provided to the system.

Similarly, if the user wants to specify a non-adaptive integration process, she can provide a script as follows:

```
ml:
  - script    : ./test_model.py
  - condition : d < 0.1 +/- 0.01
  - reliability: 0.9999
  - mode      : fp-free
  - adaptivity : none -> xx@abc.com
  - steps     : 32
```

It accepts each commit but sends the test result to the email address `xx@abc.com` after each commit. The assumption is that the developer does not have access to this email account and therefore, cannot adapt her next model.

Discussion and Future Extensions The current syntax of `ease.ml/ci` is able to capture many use cases that our users find useful in their own development process, including to reason about the accuracy difference between the new and old models, and to reason about the amount of changes in predictions between the new and old models in the test dataset. In principle, `ease.ml/ci` can support a richer syntax. We list some limitations of the current syntax that we believe are interesting directions for future work.

1. Beyond accuracy: There are other important quality metrics for machine learning that the current system does not support, e.g., F1-score, AUC score, etc. It is possible to extend the current system to accommodate these scores by replacing the Bennett's inequality with the McDiarmid's inequality, together with the sensitivity of F1-score and AUC score. In this new context, more optimizations, such as using stratified samples, are possible for skewed cases.
2. Ratio statistics: The current syntax of `ease.ml/ci` intentionally leaves out division ("/") and it would be useful for a future version to enable *relative* comparison of qualities (e.g., accuracy, F1-score, etc.).
3. Order statistics: Some users think that order statistics are also useful, e.g., to make sure the new model is among top-5 models in the development history.

Another limitation of the current system is the lack of being able to detect a domain drift or concept ship. In principle, this could be thought of as a similar process of CI – instead of fixing the test set and testing multiple models, monitoring concept shift is to fix a single model and test its generalization over multiple test sets overtime.

The current version of `ease.ml/ci` does not provide support for all these features. However, we believe that many of them can be supported by developing similar statistical techniques (see Sections 3 and 4).

2.3 System Utilities

In traditional continuous integration, the system often assumes that the user has the knowledge and competency to build the test suite all by herself. This assumption is too strong for `ease.ml/ci`— among the current users of `ease.ml/ci`, we observe that even experienced software engineers in large tech companies can be clueless on how to develop a proper testset for a given reliability requirement. One prominent contribution of `ease.ml/ci` is a collection of techniques that provide practical, but rigorous, guidelines for the user to manage testsets: *How large does the testset need to be? When does the system need a new freshly generated testset? When can the system release the*

testset and "downgrade" it into a development set? While most of these questions can be answered by experts based on heuristics and intuition, the goal of `ease.ml/ci` is to provide systematic, principled guidelines. To achieve this goal, `ease.ml/ci` provides two utilities that are not provided in systems such as Travis CI.

Sample Size Estimator This is a program that takes as input a `ease.ml/ci` script, and outputs the number of examples that the user needs to provide in the testset.

New Testset Alarm This subsystem is a program that takes as input a `ease.ml/ci` script as well as the commit history of machine learning models, and produces an alarm (e.g., by sending an email) to the user when the current testset has been used too many times and thus cannot be used to test the next committed model. Upon receiving the alarm, the user needs to provide a new testset to the system and can also release the old testset to the developer.

An impractical implementation of these two utilities is easy — the system alarms the user to request a new testset after every commit and estimates the testset size using the Hoeffding bound. However, this can result in testsets that require tremendous labeling effort, which is not always feasible.

What is "Practical?" The practicality is certainly user dependent. Nonetheless, from our experience working with different users, we observe that providing $30,000$ to $60,000$ labels for every 32 model evaluations seems reasonable for many users: $30,000$ to $60,000$ is what 2 to 4 engineers can label in a day (8 hours) at a rate of 2 seconds per label, and 32 model evaluations imply (on average) one commit per day in a month. Under this assumption, the user only needs to spend one day per month to provide test labels with a reasonable number of labelers. If the user is not able to provide this amount of labels, a "cheap mode", where the number of labels per day is easily reduced by a factor 10x, is achieved for most of the common conditions by increasing the error tolerance by a single or two percentage points.

Therefore, to make `ease.ml/ci` a useful tool for real-world users, these utilities need to be implemented in a more practical way. The technical contribution of `ease.ml/ci` is a set of techniques that we will present next, which can reduce the number of samples the system requests from the user by up to two orders of magnitude.

3 BASELINE IMPLEMENTATION

We describe the techniques to implement `ease.ml/ci` for user-specified conditions in the most general case. The techniques that we use involve standard Hoeffding inequality and a technique similar to Ladder (Blum & Hardt, 2015) in the adaptive case. This implementation is general enough to support all user-specified conditions currently supported in `ease.ml/ci`, however, it can be made more practical when the test conditions satisfy certain conditions. We leave optimizations for specific conditions to Section 4.

3.1 Sample Size Estimator for a Single Model

Estimator for a Single Variable One building block of `ease.ml/ci` is the estimator of the number of samples one needs to estimate one variable (n, o, and d) to ϵ accuracy with $1 - \delta$ probability. We construct this estimator using the standard Hoeffding bound.

A sample size estimator $n : \mathcal{V} \times [0,1]^3 \mapsto \mathbb{N}$ is a function that takes as input a variable, its dynamic range, error tolerance and success rate, and outputs the number of samples one needs in a testset. With the standard Hoeffding bound,

$$n(v, r_v, \epsilon, \delta) = \frac{-r_v^2 \ln \delta}{2\epsilon^2}$$

where r_v is the dynamic range of the variable v, ϵ the error tolerance, and $1 - \delta$ the success probability.

Recall that we makes use of the exact grammar used to define the test conditions. A formal definition of the syntax can be found in Appendix A.1.

Estimator for a Single Clause Given a clause C (e.g. $n - o > 0.01$) with a left-hand side expression Φ, a comparison operator `cmp` ($>$ or $<$), and a right-hand side constant, the sample size estimator returns the number of samples one needs to provide an (ϵ, δ)-estimation of the left-hand side expression. This can be done with a trivial recursion:

1. $n(\texttt{EXP = c * v}, \epsilon, \delta) = n(v, r_v, \epsilon/c, \delta)$, where c is a constant. We have $n(\texttt{c * v}, \epsilon, \delta) = \frac{-c^2 r_v^2 \ln \delta}{2\epsilon^2}$.

2. $n(\texttt{EXP1 + EXP2}, \epsilon, \delta) = \max\{n(\texttt{EXP1}, \epsilon_1, \frac{\delta}{2}), n(\texttt{EXP2}, \epsilon_2, \frac{\delta}{2})\}$, where $\epsilon_1 + \epsilon_2 < \epsilon$. The same equality holds similarly for $n(\texttt{EXP1 - EXP2}, \epsilon, \delta)$.

Estimator for a Single Formula Given a formula F that is a conjunction over k clauses $C_1, ..., C_k$, the sample size estimator needs to guarantee that it can satisfy each of the clause C_i. One way to build such an estimator is

3. $n(\texttt{F} = C_1 \wedge \ldots \wedge C_k, \epsilon, \delta) = \max_i n(C_i, \epsilon, \frac{\delta}{k})$.

Example Given a formula F, we now have a simple algorithm for sample size estimation. For

```
F :- n - 1.1 * o > 0.01 +/- 0.01 /\ d < 0.1 +/- 0.01
```

the system solves an optimization problem:

$$n(\texttt{F}, \epsilon, \delta) = \min_{\substack{\epsilon_1 + \epsilon_2 = \epsilon \\ \epsilon_1, \epsilon_2 \in [0,1]}} \max\{\frac{-\ln \frac{\delta}{4}}{2\epsilon_1^2}, \frac{-1.1^2 \ln \frac{\delta}{4}}{2\epsilon_2^2}, \frac{-\ln \frac{\delta}{2}}{2\epsilon^2}\}.$$

3.2 Non-Adaptive Scenarios

In the non-adaptive scenario, the system evaluates H models, without releasing the result to the developer. The result can be released to the user (the integration team).

Sample Size Estimation Estimation of sample size is easy in this case because all H models are independent. With probability $1 - \delta$, ease.ml/ci returns the right answer for each of the H models, the number of samples one needs for formula F is simply $n(\text{F}, \epsilon, \frac{\delta}{H})$. This follows from the standard union bound. Given the number of models that user hopes to evaluate (specified in the steps field of a ease.ml/ci script), the system can then return the number of samples in the testset.

New Testset Alarm The alarm for users to provide a new testset is easy to implement in the non-adaptive scenario. The system maintains a counter of how many times the testset has been used. When this counter reaches the pre-defined budget (i.e., steps), the system requests a new testset from the user. In the meantime, the old testset can be released to the developer for future development process.

3.3 Fully-Adaptive Scenarios

In the fully-adaptive scenario, the system releases the test result (a single bit indicating pass/fail) to the developer. Because this bit leaks information from the testset to the developer, one cannot use union bound anymore as in the non-adaptive scenario.

A trivial strategy exists for such a case — for every model, uses a different testset. In this case, the number of samples required is $H \cdot n(\text{F}, \epsilon, \frac{\delta}{H})$. This can be improved by applying a adaptive argument similar to Ladder (Blum & Hardt, 2015) as follows.

Sample Size Estimation For the fully adaptive scenario, ease.ml/ci uses the following way to estimate the sample size for an H-step process. The intuition is simple. Assume that a developer is deterministic or pseudo-random, her decision on the next model only relies on all the previous pass/fail signals and the initial model H_0. For H steps, there are only 2^H possible configurations of the past pass/fail signals. As a result, one only needs to enforce the union bound on all these 2^H possibilities. Therefore, the number of samples one needs is $n(\text{F}, \epsilon, \frac{\delta}{2^H})$.

Is the Exponential Term too Impractical? The improved sample size $n(\text{F}, \epsilon, \frac{\delta}{2^H})$ is much smaller than the one, $H \cdot n(\text{F}, \epsilon, \frac{\delta}{2^H})$, required by the trivial strategy. Readers might worry about the dependency on H for the fully adaptive scenario. However, for H that is not too large, e.g., $H = 32$, the above bound can still lead to practical number of samples as the $\frac{\delta}{2^H}$ is within a logarithm term. As an example, consider the following simple condition:

```
F :- n > 0.8 +/- 0.05.
```

With $H = 32$, we have

$$n(\text{F}, \epsilon, \frac{\delta}{2^H}) = \frac{\ln 2^H - \ln \delta}{2\epsilon^2}.$$

Take $\delta = 0.0001$ and $\epsilon = 0.05$, we have $n(\text{F}, \epsilon, \frac{\delta}{2^H}) = 6,279$. Assuming the developer checks in the best model everyday, this means that every month the user needs to provide only fewer than seven thousand test samples, a requirement that is not too crazy. However, if $\epsilon = 0.01$, this blows up to $156,955$, which is less practical. We will show how to tighten this bound in Section 4 for a sub-family of test conditions.

New Testset Alarm Similar to the non-adaptive scenario, the alarm for requesting a new testset is trivial to implement — the system requests a new testset when it reaches the pre-defined budget. At that point, the system can release the testset to the developer for future development.

3.4 Hybrid Scenarios

One can obtain a better bound on the number of required samples by constraining the information being released to the developer. Consider the following scenario:

1. If a commit fails, returns Fail to the developer;
2. If a commit passes, (1) returns Pass to the developer, and (2) triggers the new testset alarm to request a new testset from the user.

Compared with the fully adaptive scenario, in this scenario, the user provides a new testset immediately after the developer commits a model that passes the test.

Sample Size Estimation Let H be the maximum number of steps the system supports. Because the system will request a new testset immediately after a model passes the test, it is not really adaptive: As long as the developer continues to use the same testset, she can assume that the last model always fails. Assume that the user is a deterministic function that returns a new model given the past history and past feedback (a stream of Fail), there are only H possible states that we need to apply union bound. This gives us the same bound as the non-adaptive scenario: $n(\text{F}, \epsilon, \frac{\delta}{H})$.

New Testset Alarm Unlike the previous two scenarios, the system will alarm the user whenever the model that she provides passes the test or reaches the pre-defined budget H, whichever comes earlier.

Discussion It might be counter-intuitive that the hybrid scenario, which leaks information to the developer, has the same sample size estimator as the non-adaptive case. Given the maximum number of steps that the testset supports, H, the hybrid scenario cannot always finish all H steps as it might require a new testset in $H' \ll H$ steps. In other words, in contrast to the fully adaptive scenario, the hybrid scenario accommodates the leaking of information not by adding more samples, but by decreasing the number of steps that a testset can support.

The hybrid scenario is useful when the test is hard to pass or fail. For example, imagine the following condition:

```
F :- n - o > 0.1 +/- 0.01
```

That is, the system only accepts commits that increase the accuracy by 10 accuracy points. In this case, the developer might take many developing iterations to get a model that actually satisfies the condition.

3.5 Evaluation of a Condition

Given a testset that satisfies the number of samples given by the sample size estimator, we obtain the estimates of the three variables used in a clause, i.e., $\hat{n}$, $\hat{o}$, and $\hat{d}$. Simply using these estimates to evaluate a condition might cause both false positives and false negatives. In `ease.ml/ci`, we instead replace the point estimates by their corresponding confidence intervals, and define a simple algebra over intervals (e.g., $[a, b] + [c, d] = [a + c, b + d]$), which is used to evaluate the left-hand side of a single clause. A clause still evaluates to $\{$`True`, `False`, `Unknown`$\}$. The system then maps this three-value logic into a two-value logic given user's choice of either `fp-free` or `fn-free`.

3.6 Use Cases and Practicality Analysis

The baseline implementation of `ease.ml/ci` relies on standard concentration bounds with simple, but novel, twists to the specific use cases. Despite its simplicity, this implementation can support real-world scenarios that many of our users find useful. We summarize five use cases and analyze the number of samples required from the user. These use cases are summarized from observing the requirements from the set of users we have been supporting over the last two years, ranging from scientists at multiple universities, to real production applications provided by high-tech companies. (`[c]` and `[epsilon]` are placeholders for constants.)

(F1: Lower Bound Worst Case Quality)

```
F1          :- n > [c] +/- [epsilon]
adaptivity :- none
mode        :- fn-free
```

This condition is used for quality control to avoid the cases that the developer accidentally commits a model that has an unacceptably low quality or has obvious quality bugs. We see many use cases of this condition in non-adaptive scenario, most of which need to be false-negative free.

(F2: Incremental Quality Improvement)

```
F2          :- n - o > [c] +/- [epsilon]
adaptivity :- full
mode        :- fp-free
([c] is small)
```

This condition is used for making sure that the machine learning application monotonically improves over time. This is important when the machine learning application is end-user facing, in which it is unacceptable for the quality to drop. In this scenario, it makes sense for the whole process to be fully adaptive and false-positive free.

$1-\delta$	ϵ	F1, F4		F2, F3	
		none	full	none	full
0.99	0.1	404	1340	1753	5496
0.99	0.05	1615	5358	7012	21984
0.99	0.025	6457	21429	28045	87933
0.99	0.01	40355	133930	175282	549581
0.999	0.1	519	1455	2214	5957
0.999	0.05	2075	5818	8854	23826
0.999	0.025	8299	23271	35414	95302
0.999	0.01	51868	145443	221333	595633
0.9999	0.1	634	1570	2674	6417
0.9999	0.05	2536	6279	10696	25668
0.9999	0.025	10141	25113	42782	102670
0.9999	0.01	63381	156956	267385	641684
0.99999	0.1	749	1685	3135	6878
0.99999	0.05	2996	6739	12538	27510
0.99999	0.025	11983	26955	50150	110038
0.99999	0.01	74894	168469	313437	687736

Figure 2. Number of samples required by different conditions, $H = 32$ steps. Red font indicates "impractical" number of samples (see discussion on practicality in Section 2.3).

(F3: Significant Quality Milestones)

```
F3          :- n - o > [c] +/- [epsilon]
adaptivity :- firstChange
mode        :- fp-free
([c] is large)
```

This condition is used for making sure that the repository only contains significant quality milestones (e.g., log models after 10 points of accuracy jump). Although the condition is syntactically the same as `F2`, it makes sense for the whole process to be hybrid adaptive and false-positive free.

(F4: No Significant Changes)

```
F4          :- d < [c] +/- [epsilon]
adaptivity :- full | none
mode        :- fn-free
([c] is large)
```

This condition is used for safety concerns similar to `F1`. When the machine learning application is end-user facing or part of a larger application, it is important that its prediction will not change significantly between two subsequent versions. Here, the process needs to be false-negative free. Meanwhile, we see use cases for both fully adaptive and non-adative scenarios.

(F5: Compositional Conditions)

```
F5 :- F4 /\ F2
```

One of the most popular test conditions is a conjunction of two conditions, `F4` and `F2`: The integration team wants to use `F4` and `F2` together so that the end-user facing application will not experience dramatic quality change.

Practicality Analysis How practical is it for our baseline implementation to support these conditions, and in which case that the baseline implementation becomes impractical?

When is the Baseline Implementation Practical? The baseline implementation, in spite of its simplicity, is practical in many cases. Figure 2 illustrates the number of samples the system requires for $H = 32$ steps. We see that, for both

`F1` and `F4`, all adaptive strategies are practical up to 2.5 accuracy points, while for `F2` and `F3`, the non-adaptive and hybrid adaptive strategies are practical up to 2.5 accuracy points and the fully adaptive strategy is only practical up to 5 accuracy points. As we see from this example, even with a simple implementation, *enforcing a rigorous guarantee for CI of machine learning is not always expensive!*

When is the Baseline Implementation Not Practical? We can see from Figure 2 the strong dependency on ϵ. This is expected because of the $O(1/\epsilon^2)$ term in the Hoeffding inequality. As a result, none of the adaptive strategy is practical up to 1 accuracy point, a level of tolerance that is important for many task-critical applications of machine learning. It is also not surprising that the fully adaptive strategy requires more samples than the non-adaptive one, and therefore becomes impractical with higher error tolerance.

4 OPTIMIZATIONS

As we see from the previous sections, the baseline implementation of `ease.ml/ci` fails to provide a practical approach for low error tolerance and/or fully adaptive cases. In this section, we describe optimizations that allow us to further improve the sample size estimator.

High-level Intuition All of our proposed techniques in this section are based on the same intuition: Tightening the sample size estimator in the worst case is hard to get better than $O(1/\epsilon^2)$; instead, we take the classic system way of thinking — *improve the the sample size estimator for a sub-family of popular test conditions*. Accordingly, `ease.ml/ci` applies different optimization techniques for test conditions of different forms.

Technical Observation 1 The intuition behind a tighter sample size estimator relies on standard techniques of tightening Hoeffding's inequality for variables with small variance. Specifically, when the new model and the old model is only different on up to $(100 \times p)\%$ of the predictions, which could be part of the test condition anyway, for data point i, the random variable $n_i - o_i$ has small variance: $\mathbb{E}\left[(n_i - o_i)^2\right] < p$, where n_i and o_i are the predictions of the new and old models on the data point i. This allows us to apply the standard Bennett's inequality.

Proposition 1 (Bennett's inequality). *Let $X_1, ..., X_n$ be independent and square integrable random variables such that for some nonnegative constant b, $|X_i| \le b$ almost surely for all $i < n$. We have*

$$\Pr\left[\left|\frac{\sum_i X_i - \mathbb{E}[X_i]}{n}\right| > \epsilon\right] \le 2\exp\left(-\frac{v}{b^2}h\left(\frac{nb\epsilon}{v}\right)\right),$$

where $v = \sum_i \mathbb{E}\left[X_i^2\right]$ and $h(u) = (1+u)\ln(1+u) - u$ for all positive u.

Technical Observation 2 The second technical observation is that, to estimate the difference of predictions between the new model and the old model, one does not need to have labels. Instead, a sample from the unlabeled dataset is enough to estimate the difference. Moreover, to estimate $n - o$ when only 10% data points have different predictions, one only needs to provide labels to 10% of the whole testset.

4.1 Pattern 1: Difference-based Optimization
The first pattern that `ease.ml/ci` searches in a formula is whether it is of the following form

```
d < A +/- B /\ n - o > C +/- D
```

which constrains the amount of changes that a new model is allowed to have while ensuring that the new model is no worse than the old model. These two clauses popularly appear in test conditions from our users: For production-level systems, developers start from an already good enough, deployed model, and spend most of their time *fine-tuning* a machine learning model. As a result, the continuous integration test must have an error tolerance as low as a single accuracy point. On the other hand, the new model will not be different from the old model significantly, otherwise more engaged debugging and investigations are almost inevitable.

Assumption. One assumption of this optimization is that it is relatively cheap to obtain unlabeled data samples, whereas it is expensive to provide labels. This is true in many of the applications. When this assumption is valid, both optimizations in Section 4.1.1 and Section 4.1.2 can be applied to this pattern; otherwise, both optimizations still apply but will lead to improvement over only a subset.

4.1.1 Hierarchical Testing
The first optimization is to test the rest of the clauses conditioned on `d < A +/- B`, which leads to an algorithm with two-level tests. The first level tests whether the difference between the new model and the old model is small enough, whereas the second level tests $(n - o)$.

The algorithm runs in two steps:

1. **(Filter)** Get an $(\epsilon', \frac{\delta}{2})$-estimator $\hat{d}$ with n' samples. Test whether $\hat{d} > A + \epsilon'$: If so, returns `False`;

2. **(Test)** Test F as in the baseline implementation (with $1 - \frac{\delta}{2}$ probability), conditioned on $d < A + 2\epsilon'$.

It is not hard to see why the above algorithm works — the first step only requires unlabeled data points and does not need human intervention. In the second step, conditioned on $d < p$, we know that $\mathbb{E}\left[(n_i - o_i)^2\right] < p$ for each data point. Combined with $|n_i - o_i| < 1$, applying Bennett's inequality we have $\Pr\left[\left|\widehat{n - o} - (n - o)\right| > \epsilon\right] \le 2\exp(-nph\left(\frac{\epsilon}{p}\right))$.

As a result, the second step needs a sample size (for non-adaptive scenario) of

$$n = \frac{\ln H - \ln \frac{\delta}{4}}{ph\left(\frac{\epsilon}{p}\right)}.$$

When $p = 0.1, 1 - \delta = 0.9999, d < 0.1$, we only need 29K samples for 32 non-adaptive steps and 67K samples for 32 fully-adaptive steps to reach an error tolerance of a single accuracy point — $10\times$ fewer than the baseline (Figure 2).

4.1.2 Active Labeling

The previous example gives the user a way to conduct 32 fully-adaptive fine-tuning steps with only 67K samples. Assume that the developer performs one commit per day, this means that we require 67K samples per month to support the continuous integration service.

One potential challenge for this strategy is that all 67K samples need to be labeled before the continuous integration service can start working. This is sometimes a strong assumption that many users find problematic. In the ideal case, we hope to interleave the development effort with the labeling effort, and amortize the labeling effort over time.

The second technique our system uses relies on the observation that, to estimate $(n - o)$, only the data points that have a different prediction between the new and old models need to be labeled. When we know that the new model predictions are only different from the old model by 10%, we only need to label 10% of all data points. It is easy to see that, every time when the developer commits a new model, we only need to provide

$$n = \frac{-\ln \frac{\delta}{4}}{ph\left(\frac{\epsilon}{p}\right)} \times p$$

labels. When $p = 0.1$ and $1 - \delta = 0.9999$, then $n = 2188$ for an error tolerance of a single accuracy point. If the developer commits one model per day, the labeling team only needs to label 2,188 samples the next day. Given a well designed interface that enables a labeling throughput of 5 seconds per label, the labeling team only needs to commit 3 hours a day! For a team with multiple engineers, this overhead is often acceptable, considering the guarantee provided by the system down to a single accuracy point.

Notice that active labeling assumes a stationary underlying distribution. One way to enforce this in the system is to ask the user to provide a pool of unlabeled data points at the same time, and then only ask for labels when needed. In this way, we do not need to draw new samples over time.

4.2 Pattern 2: Implicit Variance Bound

In many cases, the user does not provide an explicit constraint on the difference between a new model and an old model. However, many machine learning models are not

so different in their predictions. Take AlexNet, ResNet, GoogLeNet, AlexNet (Batch Normalized), and VGG for example: When applied to the ImageNet testset, these five models, developed by the ML community since 2012, only produce up to 25% different answers for top-1 *correctness* and 15% different answers for top-5 *correctness*! For a typical workload of continuous integration, it is therefore not unreasonable to expect many of the consecutive commits would have smaller difference than these ImageNet winners involving years of development.

Motivated by this observation, `ease.ml/ci` will automatically match with the following pattern

```
n - o > C +/- D.
```

When the unlabeled testset is cheap to get, the system will use one testset to estimate d up to $\epsilon = 2D$: For binary classification task, the system can use an unlabeled testset; for multi-class tasks, one can either test the *difference of predictions* on an unlabeled testset or *difference of correctness* on a labeled testset. This gives us an upper bound of $n - o$. The system then tests $n - o$ up to $\epsilon = D$ on *another* testset (different from the one used to test d). When this upper bound is small enough, the system will trigger similar optimization as in `Pattern 1`. Note that the first testset will be $16\times$ smaller than testing $n - o$ directly up to $\epsilon = D$ — $4\times$ due to a higher error tolerance, and $4\times$ due to that d has $2\times$ smaller range than $n - o$.

One caveat of this approach is that the system does not know how large the second testset would be before execution. The system uses a technique similar to active labeling by incrementally growing the labeled testset every time when a new model is committed, if necessary. Specifically, we optimize for test conditions following the pattern

```
n > A +/- B,
```

when `A` is large (e.g., 0.9 or 0.95). This can be done by first having a coarse estimation of the lower bound of n, and then conducting a finer-grained estimation conditioned on this lower bound. Note that this can only introduce improvement when the lower bound is large (e.g., 0.9).

4.3 Tight Numerical Bounds

Following (Langford, 2005), having a test condition consisting of n i.i.d random variables drawn from a Bernoulli distribution, one can simply derive a tight bound on the number of samples required to reach a (ϵ, δ) accuracy. The calculation of number of samples require the probability mass function of the Binomial distribution (sum of i.i.d Bernoulli variables). Tight bound are solved by taking the minimum of number of samples n needed, over the max unknown true mean p. This technique can also be extended to more complex queries, where the binomial distribution has to be replaced by a multimodal distribution. The exact analysis has, as for the simple case, no closed-form solution, and deriving efficient approximations is left as further work.

(a) $1-\delta=0.99$ (b) $1-\delta=0.999$

Figure 3. Comparison of Sample Size Estimators in the Baseline Implementation and the Optimized Implementation.

5 EXPERIMENTS

We focus on empirically validating the derived bounds and show `ease.ml/ci` in action next.

5.1 Sample Size Estimator

One key technique most of our optimizations relied on is that, by knowing an upper bound of the sample variance, we are able to achieve a tighter bound than simply applying the Hoeffding bound. This upper bound can either be achieved by using unlabeled data points to estimate the difference between the new and old models, or by using labeled data points but conducting a coarse estimation first. We now validate our theoretical bound and its impact on improving the label complexity.

Figure 3 illustrates the estimated error and the empirical error by *assuming* different upper bounds p, for a model with accuracy around 98%. We run GoogLeNet (Jia et al., 2014) on the infinite MNIST dataset (Bottou, 2016) and estimate the true accuracy c. Assuming a non-adaptive scenario, we obtain a range of accuracies achieved by randomly taking n data points. We then estimate the interval ϵ with the given number of samples n and probability $1-\delta$. We see that, both the baseline implementation and `ease.ml/ci` dominate the empirical error, as expected, while `ease.ml/ci` uses significantly fewer samples.[1]

Figure 4 illustrates the impact of this upper bound on improving the label complexity. We see that, the improvement increases significantly when p is reasonably small — when $p=0.1$, we can achieve almost $10\times$ improvement on the label complexity. Active labeling further increases the improvement, as expected, by another $10\times$.

5.2 `ease.ml/ci` in Action

We showcase three different test conditions for a real-world incremental development of machine learning models submitted to the SemEval-2019 Task 3 competition. The goal is to classify the emotion of the user utterance as one of the

---

[1]The empirical error was determined by taking different testsets (with the sample sample size) and measuring the gap between the δ and $1-\delta$ quantiles over the observed testing accuracies.


(a) $1-\delta=0.99$ (b) $1-\delta=0.99999$

(a) $\epsilon=0.1$ (b) $\epsilon=0.01$

Figure 4. Impact of ϵ, δ, and p on the Label Complexity.

following classes: Happy, Sad, Angry or Others.[2] The eight models developed in an incremental fashion, and submitted in that exact order to the competition (finally reaching rank 29/165) are made available together with a corresponding description of each iteration via a public repository.[3] The test data, consisting of 5,509 items was published by the organizers of the competition after its termination. This represents a non-adaptive scenario, where the developer does not get any direct feedback whilst submitting new models.

Figure 5 illustrates three similar, but different test conditions, which are implemented in `ease.ml/ci`. The first two conditions check whether the new model is better than the old one by at least 2 percentage points in a non-adaptive matter. The developer will therefore not get any direct feedback as it was the case during the competition. While query (I) does reject false positive, condition (II) does accept false negative. The third condition mimics the scenario where the user would get feedback after every commit without any false negative. All three queries were optimized by `ease.ml/ci` using Pattern 2 and exploiting the fact that between any two submission there is no more than 10% difference in prediction.

Simply using Hoeffding's inequality does not lead to a practical solution — for $\epsilon=0.02$ and $\delta=0.002$, in $H=7$ non-adaptive steps, one would need

$$n > \frac{r_v^2(\ln H - \ln \frac{\delta}{2})}{2\epsilon^2} = 44,268$$

samples. This number even grows to up to 58K in the fully adaptive case!
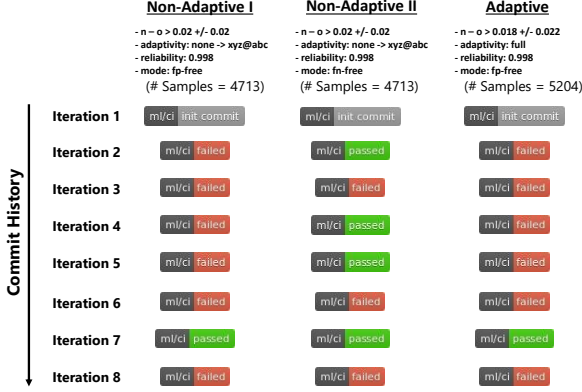
---

[2]Competition website: https://www.humanizing-ai.com/emocontext.html

[3]Github repository: https://github.com/zhaopku/ds3-emoContext

Figure 5. Continuous Integration Steps in `ease.ml/ci`.

Figure 6. Evolution of Development and Test Accuracy.

All the queries can be supported rigorously with the 5.5K test samples provided after the competition. The first two conditions can be answered within two percentage point error tolerance and 0.998 reliability. The full-adaptive query in the third scenario can only achieve a 2.2 percentage point error tolerance, as the number of labels needed would be more than 6K, with the same error tolerance as in the first two queries.

We see that, in all three scenarios, `ease.ml/ci` returns `pass`/`fail` signals that make intuitive sense. If we look at the evolution of the development and test accuracy over the eight iterations (see Figure 6, the developer would ideally want `ease.ml/ci` to accept her last commit, whereas all three queries will have the second last model chosen to be active, which correlates with the test accuracy evolution.

6 RELATED WORK

Continuous integration is a popular concept in software engineering (Duvall et al., 2007). Nowadays, it is one of the best practices that most, if not all, industrial development efforts follow. The emerging requirement of a CI engine for ML has been discussed informally in multiple blog posts and forum discussions (Lara, 2017; Tran, 2017; Stojnic, 2018a; Lara, 2018; Stojnic, 2018b). However, none of these discussions produce any rigorous solutions to testing the quality of a machine learning model, which arguably is the most important aspect of a CI engine for ML. This paper is motivated by the success of CI in industry, and aims for building the first prototype system for rigorous integration of machine learning models.

The baseline implementation of `ease.ml/ci` builds on intensive previous work on generalization and adaptive analysis. The non-adaptive version of the system is based on simple concentration inequalities (Boucheron et al., 2013) and the fully adaptive version of the system is inspired by Ladder (Blum & Hardt, 2015). Comparing to the second, `ease.ml/ci` is less restrictive on the feedback and more expressive given the specification of the test conditions. This leads to a higher number of test samples needed in general. It is well-known that the $O(1/\epsilon^2)$ sample com-

plexity of Hoeffding's inequality becomes $O(1/\epsilon)$ when the variance of the random variable σ^2 is of the same order of ϵ (Boucheron et al., 2013). In this paper, we develop techniques to adapt the same observation to a real-world scenario (Pattern 1). The technique of only labeling the difference between models is inspired by disagreement-based active learning (Hanneke et al., 2014), which illustrates the potential of taking advantage of the overlapping structure between models to decrease labeling complexity. In fact, the technique we develop implies that one can achieve $O(1/\epsilon)$ label complexity when the overlapping ratio between two models $p = O(\sqrt{\epsilon})$.

The key difference between `ease.ml/ci` and a differential privacy approach (Dwork et al., 2014) for answering statistical queries lies in the optimization techniques we design. By knowing the structure of the queries we are able to considerably lower the number of samples needed.

Conceptually, this work is inspired by the seminal series of work by Langford and others (Langford, 2005; Kääriäinen & Langford, 2005) that illustrates the possibility for generalization bound to be practically tight. The goal of this work is to build a practical system to guide the user in employing complicated statistical inequalities and techniques to achieve *practical* label complexity.

7 CONCLUSION

We have presented `ease.ml/ci`, a continuous integration system for machine learning. It provides a declarative scripting language that allows users to state a rich class of test conditions with rigorous probabilistic guarantees. We have also studied the novel practicality problem in terms of labeling effort that is specific to testing machine learning models. Our techniques can reduce the amount of required testing samples by up to two orders of magnitude. We have validated the soundness of our techniques, and showcased their applications in real-world scenarios.

REFERENCES

Blum, A. and Hardt, M. The ladder: A reliable leaderboard for machine learning competitions. In *International Conference on Machine Learning*, pp. 1006–1014, 2015.

Bottou, L. The infinite MNIST dataset. https://leon.bottou.org/projects/infimnist, February 2016.

Boucheron, S., Lugosi, G., and Massart, P. *Concentration inequalities: A nonasymptotic theory of independence*. Oxford university press, 2013.

Duvall, P. M., Matyas, S., and Glover, A. *Continuous integration: improving software quality and reducing risk*. Pearson Education, 2007.

Dwork, C., Roth, A., et al. The algorithmic foundations of differential privacy. *Foundations and Trends® in Theoretical Computer Science*, 9(3–4):211–407, 2014.

Dwork, C., Feldman, V., Hardt, M., Pitassi, T., Reingold, O., and Roth, A. The reusable holdout: Preserving validity in adaptive data analysis. *Science*, 349(6248):636–638, 2015.

Hanneke, S. et al. Theory of disagreement-based active learning. *Foundations and Trends® in Machine Learning*, 7(2-3):131–309, 2014.

Jia, Y., Shelhamer, E., Donahue, J., Karayev, S., Long, J., Girshick, R., Guadarrama, S., and Darrell, T. Caffe: Convolutional architecture for fast feature embedding. *arXiv preprint arXiv:1408.5093*, 2014.

Kääriäinen, M. and Langford, J. A comparison of tight generalization error bounds. In *Proceedings of the 22nd international conference on Machine learning*, pp. 409–416. ACM, 2005.

Langford, J. Tutorial on practical prediction theory for classification. *Journal of machine learning research*, 6 (Mar):273–306, 2005.

Lara, A. F. Continuous integration for ml projects. https://medium.com/onfido-tech/continuous-integration-for-ml-projects-e11bc1a4d34f, October 2017.

Lara, A. F. Continuous delivery for ml models. https://medium.com/onfido-tech/continuous-delivery-for-ml-models-c1f9283aa971, July 2018.

Stojnic, R. Continuous integration for machine learning. https://medium.com/@rstojnic/continuous-integration-for-machine-learning-6893aa867002, April 2018a.

Stojnic, R. Continuous integration for machine learning. https://www.reddit.com/r/MachineLearning/comments/8bq5la/d_continuous_integration_for_machine_learning/, April 2018b.

Tran, D. Continuous integration for data science. http://engineering.pivotal.io/post/continuous-integration-for-data-science/, February 2017.

Van Vliet, H., Van Vliet, H., and Van Vliet, J. *Software engineering: principles and practice*, volume 13. John Wiley & Sons, 2008.

A SYNTAX AND SEMANTICS

A.1 Syntax of a Condition

To specify the condition, which will be tested by `ease.ml/ci` whenever a new model is committed, the user makes use of the following grammar:

```
c    :- floating point constant
v    :- n | o | d
op1 :- + | -
op2 :- *
EXP :- v | v op1 EXP | EXP op2 c

cmp :- > | <
C    :- EXP cmp c +/- c

F    :- C | C /\ F
```

`F` is the final condition, which is a conjunction of a set of clauses `C`. Each clause is a comparison between an expression over $\{n, o, d\}$ and a constant, with an error tolerance following the symbol `+/-`. For example, two expressions that we focus on optimizing can be specified as follows:

```
 n - o > 0.02 +/- 0.01 /\ d < 0.1 +/- 0.01
```

in which the first clause

```
        n - o > 0.02 +/- 0.01
```

requires that the new model have an accuracy that is two points higher than the old model, with an error tolerance of one point, whereas the clause

```
        d < 0.1 +/- 0.01
```

requires that the new model can only change 10% of the old predictions, with an error tolerance of 1%.

A.2 Semantics of Continuous Integration Tests

Unlike traditional continuous integration, all three variables used in `ease.ml/ci`, i.e., $\{n, o, d\}$, are *random variables*. As a result, the evaluation of an `ease.ml/ci` condition is inherently *probabilistic*. There are two additional parameters that the user needs to provide, which would define the semantics of the test condition: (1) δ, the probability with which the test process is allowed to be incorrect, which is usually chosen to be smaller than 0.001 or 0.0001 (i.e., 0.999 or 0.9999 success rate); and (2) `mode` chosen from $\{$`fp-free, fn-free`$\}$, which specifies whether the test is *false-positive free* or *false-negative free*. The semantics are, with probability $1 - \delta$, the output of `ease.ml/ci` is free of false positives or false negatives.

The notion of false positives or false negatives is related to the fundamental trade-off between the "type I" error and the "type II" error in statistical hypothesis testing. Consider

```
        x < 0.1 +/- 0.01.
```

Suppose that the real *unknown* value of x is x^*. Given an estimator $\hat{x}$, which, with probability $1 - \delta$, satisfies

$$\hat{x} \in [x^* - 0.01, x^* + 0.01],$$

what should be the testing outcome of this condition? There are three cases:

1. When $\hat{x} > 0.11$, the condition should return `False` because, given $x^* < 0.1$, the probability of having $\hat{x} > 0.11 > x^* + 0.01$ is less than δ.

2. When $\hat{x} < 0.09$, the condition should return `True` because, given $x^* > 0.1$, the probability of having $\hat{x} < 0.09 < x^* - 0.01$ is less than δ.

3. When $0.09 < \hat{x} < 0.11$, the outcome cannot be determined: Even if $\hat{x} > 0.1$, there is no way to tell whether the real value x^* is larger or smaller than 0.1. In this case, the condition evaluates to `Unknown`.

The parameter `mode` allows the system to deal with the case that the condition evaluates to `Unknown`. In the `fp-free` mode, `ease.ml/ci` treats `Unknown` as `False` (thus rejects the commit) to ensure that whenever the condition evaluates to `True` using $\hat{x}$, the same condition is always `True` for x^*. Similarly, in the `fn-free` mode, `ease.ml/ci` treats `Unknown` as `True` (thus accepts the commit). The false positive rate (resp. false negative rate) in the `fn-free` (resp. `fp-free`) mode is specified by the error tolerance.